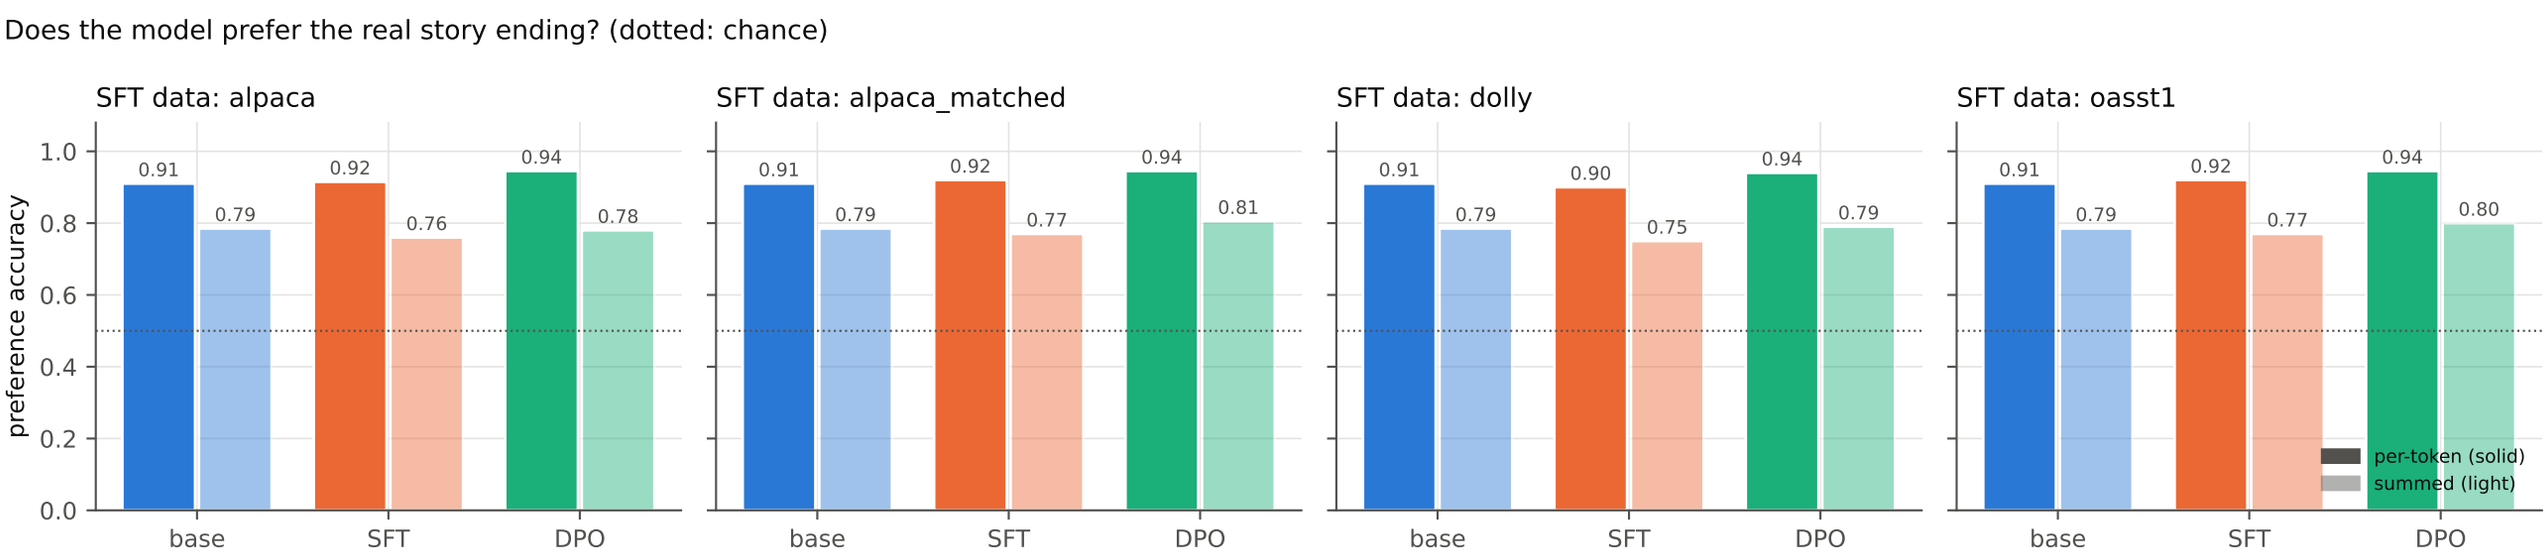
Does the model prefer the real story ending? (dotted: chance)
SFT data: alpaca
SFT data: alpaca_matched
SFT data: dolly
SFT data: oasst1
preference accuracy
1.0
0.8
0.6
0.4
0.2
0.0
0.91
0.79
0.92
0.76
0.94
0.78
0.91
0.79
0.92
0.77
0.94
0.81
0.91
0.79
0.90
0.75
0.94
0.79
0.91
0.79
0.92
0.77
0.94
0.80
base
SFT
DPO
per-token (solid)
summed (light)